

A (not so) short introduction to Reinforcement Learning

Jost Tobias Springenberg
springj@informatik.uni-freiburg.de

Machine Learning Lab - University of Freiburg

December 23, 2016

Disclaimer

- ▶ Reinforcement Learning (RL) is a huge field, I will focus on the very basics that you need for your exercise

Overview

1. What is a MDP ?
2. The Dimensions of RL
3. RL Definitions
4. Model Based RL
5. Model Free RL
 - 5.1 Q-Learning and TD methods

What is a MDP ?

- ▶ Before I detail anything complicated let us look at the framework of Markov Decision Processes (MDPs) which we will use to describe all RL algorithms
- ▶ The situation an MDP seeks to model:






Figure: Taken from RL lecture by Martin Riedmiller

What is a MDP ?




What is a MDP ?



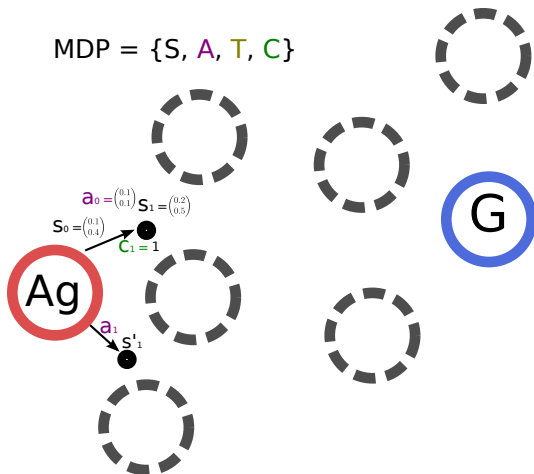
What is a MDP ?

$$\text{MDP} = \{S, A, T,$$



What is a MDP ?


$$\text{MDP} = \{S, A, T, C\}$$



What is a MDP ?


$$\text{MDP} = \{S, A, T, C\}$$

$$\pi(s) = p(a \mid s)$$



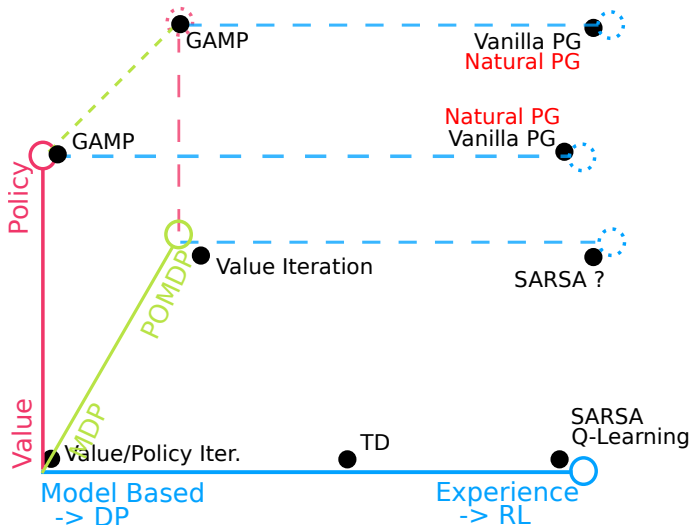
The Dimensions of RL

- ▶ We are going to look at RL methods by considering three axes:



The Dimensions of RL

- We are going to look at RL methods by considering three axes:



RL Definitions

Ok let us formalize

- ▶ A MDP is a 5-tuple

$$\text{MDP} = (S, A, P, R)$$

RL Definitions

Ok let us formalize

- ▶ A MDP is a 5-tuple

$$\text{MDP} = (S, A, P, R)$$

- ▶ Discrete timesteps t with states and actions $s_t \in S$, $a_t \in A$

RL Definitions

Ok let us formalize

- ▶ A MDP is a 5-tuple

$$\text{MDP} = (S, A, P, R)$$

- ▶ Discrete timesteps t with states and actions $s_t \in S$, $a_t \in A$
- ▶ Deterministic rewards: $r_t = r(s_t, a_t) \mapsto \mathbb{R}$

RL Definitions

Ok let us formalize

- ▶ A MDP is a 5-tuple

$$\text{MDP} = (S, A, P, R)$$

- ▶ Discrete timesteps t with states and actions $s_t \in S, a_t \in A$
- ▶ Deterministic rewards: $r_t = r(s_t, a_t) \mapsto \mathbb{R}$
- ▶ Probabilistic transition model P : $p(s_{t+1} \mid s_t, a_t)$
 - ▶ Assume the *markov property*: the future depends on the past only through the present state

RL Definitions

Ok let us formalize

- ▶ A MDP is a 5-tuple

$$\text{MDP} = (S, A, P, R)$$

- ▶ Discrete timesteps t with states and actions $s_t \in S, a_t \in A$
- ▶ Deterministic rewards: $r_t = r(s_t, a_t) \mapsto \mathbb{R}$
- ▶ Probabilistic transition model P : $p(s_{t+1} \mid s_t, a_t)$
 - ▶ Assume the *markov property*: the future depends on the past only through the present state
 - $p(s_{t+1} \mid s_t, a_t)$ independent of previous states and actions

RL Definitions

Ok let us formalize

- ▶ A MDP is a 5-tuple

$$\text{MDP} = (S, A, P, R)$$

- ▶ Discrete timesteps t with states and actions $s_t \in S, a_t \in A$
- ▶ Deterministic rewards: $r_t = r(s_t, a_t) \mapsto \mathbb{R}$
- ▶ Probabilistic transition model P : $p(s_{t+1} \mid s_t, a_t)$
 - ▶ Assume the *markov property*: the future depends on the past only through the present state
 - $p(s_{t+1} \mid s_t, a_t)$ independent of previous states and actions
- ▶ The agents policy is described as: $\pi(s_t, a_t) = p(a_t \mid s_t)$

Model Based RL

→ **Question** How can we figure out what action to take in a state ?

Model Based RL

- Question How can we figure out what action to take in a state ?
- ▶ We can assign a Value V to each state as follows.
 - ▶ Let us assume a finite horizon

$$V_T^\pi(s) = \mathbb{E}_{\pi, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_0 = s \right\}$$

Model Based RL

→ **Question** How can we figure out what action to take in a state ?

- ▶ We can assign a Value V to each state as follows.
- ▶ Let us assume a finite horizon

$$V_T^\pi(s) = \mathbb{E}_{\pi, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_0 = s \right\}$$

- ▶ **Note:** if we knew the $V^\pi \forall \pi$ behaving optimally: pick policy with maximum expected reward when starting in s_0

Model Based RL

→ **Question** How can we figure out what action to take in a state ?

- ▶ We can assign a Value V to each state as follows.
- ▶ Let us assume a finite horizon

$$V_T^\pi(s) = \mathbb{E}_{\pi, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_0 = s \right\}$$

- ▶ **Note:** if we knew the $V^\pi \forall \pi$ behaving optimally: pick policy with maximum expected reward when starting in s_0
- ▶ **Problem:** But we don't really want to enumerate all of these

Model Based RL

→ **Question** How can we figure out what action to take in a state ?

- ▶ We can assign a Value V to each state as follows.
- ▶ Let us assume a finite horizon

$$V_T^\pi(s) = \mathbb{E}_{\pi, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_0 = s \right\}$$

- ▶ **Note:** if we knew the $V^\pi \forall \pi$ behaving optimally: pick policy with maximum expected reward when starting in s_0
- ▶ **Problem:** But we don't really want to enumerate all of these
- ▶ **Solution:** Classic dynamic programming

Model Based RL - Dynamic Programming (DP)

- ▶ **Solution:** Classic dynamic programming
- ▶ **Wanted:** $\pi^* = \arg \max_{\pi} V_T^{\pi}(s)$

Model Based RL - Dynamic Programming (DP)

- ▶ **Solution:** Classic dynamic programming
- ▶ **Wanted:** $\pi^* = \arg \max_{\pi} V_T^{\pi}(s)$
- ▶ Bellmans optimality principle states that for optimal π^* we can reformulate

$$\begin{aligned} V_T^{\pi^*}(s_t) &= \mathbb{E}_{\pi^*, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_0 = s \right\} \\ &= \max_a r(s, a) + \sum_{s' \in S} p(s' \mid s, a) V_{T-1}^{\pi^*}(s') \end{aligned}$$

Model Based RL - Dynamic Programming (DP)


- ▶ **Solution:** Classic dynamic programming
- ▶ **Wanted:** $\pi^* = \arg \max_{\pi} V_T^{\pi}(s)$
- ▶ Bellmans optimality principle states that for optimal π^* we can reformulate

$$\begin{aligned} V_T^{\pi^*}(s_t) &= \mathbb{E}_{\pi^*, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_o = s \right\} \\ &= \max_a r(s, a) + \sum_{s' \in S} p(s' \mid s, a) V_{T-1}^{\pi^*}(s') \end{aligned}$$


→ Set $V_0^{\pi^*}(s') = \max_a r(s, a)$ and iterate

Model Based RL - Dynamic Programming (DP)


- But is this actually any use ? Show me a solvable problem!




Model Based RL - Dynamic Programming (DP)



Model Based RL - Dynamic Programming (DP)



Model Based RL - Dynamic Programming (DP)



→ lather, rinse, repeat

Model Based RL - DP Problems


- ▶ DP works great but only allows Agent to “live” for N steps ...

Model Based RL - DP Problems

- ▶ DP works great but only allows Agent to “live” for N steps ...
- ▶ That is not the problem we really care for at all !

Model Based RL - DP Problems

- ▶ DP works great but only allows Agent to “live” for N steps ...
- ▶ That is not the problem we really care for at all !
- ▶ Remember (continuous states/actions, infinite horizon):



Model Based RL - DP Problems

- Recall the definition of the state value:

$$V_T^\pi(s) = \mathbb{E}_{\pi, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_0 = s \right\}$$

Model Based RL - DP Problems

- ▶ Recall the definition of the state value:

$$V_T^\pi(s) = \mathbb{E}_{\pi, P} \left\{ \sum_{t=0}^{T-1} r_t \mid s_0 = s \right\}$$

- ▶ Why don't we consider an infinite horizon ?

$$V_\infty^\pi(s) = \mathbb{E}_{\pi, P} \left\{ \sum_{t=0}^{\infty} r_t \mid s_0 = s \right\}$$

Model Based RL - DP Problems

- ▶ Recall the definition of the state value:

$$V_T^\pi(s) = \mathbb{E}_{\pi,P}\left\{\sum_{t=0}^{T-1} r_t \mid s_o = s\right\}$$

- ▶ Why don't we consider an infinite horizon ?

$$V_\infty^\pi(s) = \mathbb{E}_{\pi,P}\left\{\sum_{t=0}^{\infty} r_t \mid s_o = s\right\}$$

- ▶ Not bounded, we don't want infinite costs
- Solve this via discounting with $\gamma \leq 1$

$$V^\pi(s) = \mathbb{E}_{\pi,P}\left\{\sum_{t=0}^{\infty} \gamma^t r_t \mid s_o = s\right\}$$

Model Based RL - DP Problems

- ▶ Recall the definition of the state value:

$$V_T^\pi(s) = \mathbb{E}_{\pi,P}\left\{\sum_{t=0}^{T-1} r_t \mid s_o = s\right\}$$

- ▶ Why don't we consider an infinite horizon ?

$$V_\infty^\pi(s) = \mathbb{E}_{\pi,P}\left\{\sum_{t=0}^{\infty} r_t \mid s_o = s\right\}$$

- ▶ Not bounded, we don't want infinite costs
- Solve this via discounting with $\gamma \leq 1$

$$V^\pi(s) = \mathbb{E}_{\pi,P}\left\{\sum_{t=0}^{\infty} \gamma^t r_t \mid s_o = s\right\}$$

- ▶ Discounting is like assuming a non-zero interest rate — money arriving in the future is worth less than money arriving now.

Model Based RL - Value iteration

- ▶ Can we solve for the optimal value function (and hence policy) ?

Model Based RL - Value iteration

- ▶ Can we solve for the optimal value function (and hence policy) ?
- ▶ Solution is to iteratively do DP like updates with discounting

Algorithm 2: Value Iteration

```
1 repeat
2   for  $s \in S$  do
3     update  $V$  values using current estimate  $V_i$ 
4      $V_{i+1}^*(s) \leftarrow \max_a \sum_{s' \in S} p(s' | s, a)(r(s, a) + \gamma V_i^*(s'))$ 
5 until converged
6 return  $V^*$ 
```

Model Based RL - Value iteration

- ▶ Can we solve for the optimal value function (and hence policy) ?
- ▶ Solution is to iteratively do DP like updates with discounting

Algorithm 3: Value Iteration

```
1 repeat
2   for  $s \in S$  do
3     update  $V$  values using current estimate  $V_i$ 
4      $V_{i+1}^*(s) \leftarrow \max_a \sum_{s' \in S} p(s' | s, a)(r(s, a) + \gamma V_i^*(s'))$ 
5 until converged
6 return  $V^*$ 
```

→ Converges under mild assumptions

Model Based RL - Value iteration

- ▶ Can we solve for the optimal value function (and hence policy) ?
- ▶ Solution is to iteratively do DP like updates with discounting

Algorithm 4: Value Iteration

```
1 repeat
2   for  $s \in S$  do
3     update  $V$  values using current estimate  $V_i$ 
4      $V_{i+1}^*(s) \leftarrow \max_a \sum_{s' \in S} p(s' | s, a)(r(s, a) + \gamma V_i^*(s'))$ 
5 until converged
6 return  $V^*$ 
```

- Converges under mild assumptions
- However only in the limit of infinite iterations ...

Model Based RL - Further Problems

- ▶ But what if we don't have a model ?
 - Learn a model (see e.g. PILCO, or our ADPRL paper)

movie1

- Use model free methods (e.g. Q-learning, SARSA)

Model Based RL - Further Problems

- ▶ But what if we don't have a model ?
 - Learn a model (see e.g. PILCO, or our ADPRL paper)

movie1

- Use model free methods (e.g. Q-learning, SARSA)
- ▶ And what if we have continuous states ?
 - Use function approximation (e.g. a Neural Network)

Model Free RL

- But what if we don't have a model ?

Algorithm 5: Value Iteration

```
1 repeat
2   for  $s \in S$  do
3     update  $V$  values using current estimate  $V_i$ 
4      $V_{i+1}^*(s) \leftarrow \max_a \sum_{s' \in S} p(s' | s, a) (r(s, a) + \gamma V_i^*(s'))$ 
5 until converged
6 return  $V^*$ 
```

(Somewhat) Model Free RL - TD-learning

- ▶ Somewhat Model Free RL → use sampling

Algorithm 6: TD learning

```
1 repeat
2   choose  $a = \arg \max_a r(s, a) + \sum_{s'} p(s' | s, a) V(s')$ 
3   record transition  $(s, a, r, s')$ 
4   update  $V$  using sample (td) error  $\delta_t = r(s, a) + \gamma V_i(s') - V_i(s)$ 
5    $V_{i+1}(s) \leftarrow V_i(s) + \alpha \delta_t$ 
6 until converged
7 return  $V$ 
```

(Somewhat) Model Free RL - TD-learning

- ▶ Almost there!

(Somewhat) Model Free RL - TD-learning

- ▶ Almost there!
- ▶ Need a way to query the value of actions without using the model

(Somewhat) Model Free RL - TD-learning

- ▶ Almost there!
- ▶ Need a way to query the value of actions without using the model
- ▶ Introducing the Q-function

$$Q^{\pi}(s, a) = r(s, a) + \sum_{s'} p(s' \mid s, a) V^{\pi}(s')$$

Model Free RL - Q-learning

Algorithm 7: Q-learning

```
1 repeat
2   choose  $a = \arg \max_a Q(s, a)$ 
3   record transition  $(s, a, r, s')$ 
4   update  $Q$  using sample Q-error
       $\delta_t = r(s, a) + \gamma \max_{a'} Q_i(s', a') - Q_i(s, a)$ 
5    $Q_{i+1}(s, a) \leftarrow Q_i(s, a) + \alpha \delta_t$ 
6 until converged
7 return  $Q$ 
```

Model Free RL - Q-learning

Algorithm 8: Q-learning

```
1 repeat
2   choose  $a = \begin{cases} \arg \max_a Q(s, a) & \text{with probability } \epsilon \\ a \in A & \text{else choose randomly} \end{cases}$ 
3   record transition  $(s, a, r, s')$ 
4   update  $Q$  using sample Q-error
    $\delta_t = r(s, a) + \gamma \max_{a'} Q_i(s', a') - Q_i(s, a)$ 
5    $Q_{i+1}(s, a) \leftarrow Q_i(s, a) + \alpha \delta_t$ 
6 until converged
7 return  $Q$ 
```


Model Free RL - Q-learning

→ Remarkably Q-learning converges

- ▶ Okay, we don't need a model anymore, let's see some more applications!


Model Free RL - Q-learning

- ▶ Finally solving some important problems!



Model Free RL - Q-learning

- ▶ Solving FlappyBird with a simple table based Q-learning approach (see <http://sarvagyaavaish.github.io/FlappyBirdRL/>)
- ▶ 2-dimensional State-space S : (x, y) pixel distance to tube discretized in 4 pixel steps



Model Free RL - Q-learning

- ▶ Solving FlappyBird with a simple table based Q-learning approach (see <http://sarvagyaish.github.io/FlappyBirdRL/>)

Model Free RL - Q-learning

- ▶ Solving FlappyBird with a simple table based Q-learning approach (see <http://sarvagyaisha.github.io/FlappyBirdRL/>)
- ▶ 2-dimensional State-space S : (x, y) pixel distance to tube discretized in 4 pixel steps

Model Free RL - Q-learning

- ▶ Solving FlappyBird with a simple table based Q-learning approach (see <http://sarvagyaivaish.github.io/FlappyBirdRL/>)
- ▶ 2-dimensional State-space S : (x, y) pixel distance to tube discretized in 4 pixel steps
- ▶ 2-actions $A = \{\text{tap}, \text{donothing}\}$

Model Free RL - Q-learning

- ▶ Solving FlappyBird with a simple table based Q-learning approach (see <http://sarvagyaivaish.github.io/FlappyBirdRL/>)
- ▶ 2-dimensional State-space S : (x, y) pixel distance to tube discretized in 4 pixel steps
- ▶ 2-actions $A = \{\text{tap}, \text{donothing}\}$
- ▶ Reward

$$r(s, a) = \begin{cases} 1 & \text{if alive} \\ -1000 & \text{otherwise} \end{cases} \quad (1)$$

Model Free RL - Q-learning

- ▶ Solving FlappyBird with a simple table based Q-learning approach (see <http://sarvagyaivaish.github.io/FlappyBirdRL/>)
- ▶ 2-dimensional State-space S : (x, y) pixel distance to tube discretized in 4 pixel steps
- ▶ 2-actions $A = \{\text{tap}, \text{donothing}\}$
- ▶ Reward

$$r(s, a) = \begin{cases} 1 & \text{if alive} \\ -1000 & \text{otherwise} \end{cases} \quad (1)$$

- ▶ Solved using on-line Q-learning (training took about 6-7 hours)
- ▶ Show VIDEO

Model Free RL - Q-learning

So are we there yet ?

→ Problems with Q-learning

Model Free RL - Q-learning

So are we there yet ?

→ Problems with Q-learning

→ Curse of dimensionality I: table is not going to work for high dimensions

Model Free RL - Q-learning

So are we there yet ?

→ Problems with Q-learning

→ Curse of dimensionality I: table is not going to work for high dimensions

→ We have only considered discrete states, next: function approximation

Model Free RL - Q-learning

So are we there yet ?

- Problems with Q-learning
 - Curse of dimensionality I: table is not going to work for high dimensions
 - We have only considered discrete states, next: function approximation
 - Curse of dimensionality II: Efficient Exploration is key! Even with only 2 state-dimensions + 2 actions FlappyBird took hours ..

Model Free RL - Q-learning

Modern approaches / Function approximation

- ▶ To (partially) solve the curse of dimensionality one can try to use function approximation

Model Free RL - Q-learning

Modern approaches / Function approximation

- ▶ To (partially) solve the curse of dimensionality one can try to use function approximation

→ **Idea:** replace Q-table by a function approximator

- ▶ use Q-error $\delta_t = r(s, a) + \gamma \max_a Q_i(s', a) - Q_i(s, a)$ as the gradient

Model Free RL - From Q-learning to DQN

Algorithm 9: Approximate Q-learning

```
1 repeat
2   choose  $a = \begin{cases} \arg \max_a Q(s, a) & \text{with probability } \epsilon \\ a \in A & \text{else choose randomly} \end{cases}$ 
3   record transition  $(s, a, r, s')$ 
4   update  $Q$  using sample Q-error
   
$$\delta_t = r(s, a) + \gamma \max_{a'} Q_i(s', a') - Q_i(s, a)$$

5   
$$Q_{i+1}(s, a) \leftarrow Q_i(s, a) + \alpha \delta_t$$

6 until converged
7 return  $Q$ 
```

Model Free RL - From Q-learning to DQN

Algorithm 10: Approximate Q-learning

```
1 repeat
2   choose  $a = \begin{cases} \arg \max_a Q_W(s, a) & \text{with probability } \epsilon \\ a \in A & \text{else choose randomly} \end{cases}$ 
3   record transition  $(s, a, r, s')$ 
4   update  $Q_W$  using sample Q-error
   
$$\delta_t = r(s, a) + \gamma \max_{a'} Q_W(s', a') - Q_W(s, a)$$

5   
$$Q_{i+1}(s, a) \leftarrow Q_i(s, a) + \alpha \delta_t$$

6 until converged
7 return  $Q$ 
```

Model Free RL - From Q-learning to DQN

Algorithm 11: Approximate Q-learning

1 **repeat**

2 chose $a^* = \begin{cases} \max_a Q_W(s, a) & \text{with probability } \epsilon \\ a \in A & \text{else chose randomly} \end{cases}$

3 record transition (s, a, r, s')

4 update Q_W using sample Q-error as derivative

$$\delta_t = r(s, a) + \gamma \max_{a'} Q_W(s', a') - Q_W(s, a)$$

5
$$W \leftarrow W + \alpha \left(\delta_t \frac{\partial Q_W(s, a)}{\partial W} \right)$$

6 **until** *converged*

7 **return** Q_W

Model Free RL - Q-learning

Modern approaches / Function approximation

- ▶ To (partially) solve the curse of dimensionality one can try to use function approximation
- **Idea:** replace Q-table by a function approximator
 - ▶ use Q-error $\delta_t = r(s, a) + \gamma \max_a Q_i(s', a) - Q_i(s, a)$ as the gradient
- **Problem:** What is the function we are minimizing anyways ?
Convergence ?

Model Free RL - Q-learning

Modern approaches / Function approximation

- ▶ To (partially) solve the curse of dimensionality one can try to use function approximation
- **Idea:** replace Q-table by a function approximator
 - ▶ use Q-error $\delta_t = r(s, a) + \gamma \max_a Q_i(s', a) - Q_i(s, a)$ as the gradient
- **Problem:** What is the function we are minimizing anyways ?
Convergence ?
 - ▶ Unfortunately approximate Q-learning may not converge

Model Free RL - Q-learning

Modern approaches / Function approximation

- ▶ To (partially) solve the curse of dimensionality one can try to use function approximation
- **Idea:** replace Q-table by a function approximator
 - ▶ use Q-error $\delta_t = r(s, a) + \gamma \max_a Q_i(s', a) - Q_i(s, a)$ as the gradient
- **Problem:** What is the function we are minimizing anyways ?
Convergence ?
 - ▶ Unfortunately approximate Q-learning may not converge
- **Problem:** Non-stationarity/Floating targets ? If Q changes the function we minimize changes

Model Free RL - Q-learning


Modern approaches / Function approximation

- ▶ To (partially) solve the curse of dimensionality one can try to use function approximation
- **Idea:** replace Q-table by a function approximator
 - ▶ use Q-error $\delta_t = r(s, a) + \gamma \max_a Q_i(s', a) - Q_i(s, a)$ as the gradient
- **Problem:** What is the function we are minimizing anyways ?
Convergence ?
 - ▶ Unfortunately approximate Q-learning may not converge
- **Problem:** Non-stationarity/Floating targets ? If Q changes the function we minimize changes
- Lots of exciting ongoing work: LSTD / NFQ / DQN etc.

Next exercise

NOTE: An exercise sheet is on the course webpage and in the git repository!

- 1) By hand, on a piece of paper solve the following discretized maze via Q-learning



- 2) Implement a simple neural Q-learning agent using a simple map in the same environment you already know
 - ▶ API slightly changed so that you get a reward in each step
 - ▶ hand in report showing the solution to 1) and explain what you did for 2)
 - ▶ Exercise sheet and update code available **Monday 19.12.16** deadline for exercise **23.01.17**
 - ▶ Bonus points for trying some more advanced ideas (e.g. target network, experience replay as in DQN)

Model Free RL - From Q-learning to DQN

Algorithm 12: Smoothed Approximate Q-learning (DQN)

```
1 repeat
2   chose  $a^* = \begin{cases} \max_a Q_W(s, a) & \text{with probability } \epsilon \\ a \in A & \text{else chose randomly} \end{cases}$ 
3   record transition  $(s, a, r, s')$ 
4   update  $Q_W$  using sample Q-error as derivative
    $\delta_t = r(s, a) + \gamma \max_{a'} Q_W(s', a') - Q_W(s, a)$ 
5    $W \leftarrow W - \alpha \left( \delta_t \frac{\partial Q_W(s, a)}{\partial W} \right)$ 
6 until converged
7 return  $Q_W$ 
```

Model Free RL - From Q-learning to DQN

Algorithm 13: Smoothed Approximate Q-learning (DQN)

```
1   $T = \{\}$ 
2  repeat
3      chose  $a^* = \begin{cases} \max_a Q_W(s, a) & \text{with probability } \epsilon \\ a \in A & \text{else chose randomly} \end{cases}$ 
4      record transition  $T = T \cup \{(s, a, r, s')\}$ 
5      update  $Q_W$  using sample Q-error as derivative
6      for  $s \in [1, K]$  do
7           $(s, a, r, s') \sim T$ 
8           $\delta = \delta + r(s, a) + \gamma \max_{a'} Q_W(s', a') - Q_W(s, a)$ 
9           $W \leftarrow W + \alpha \left( \delta \frac{\partial Q_W(s, a)}{\partial W} \right)$ 
10 until converged
11 return  $Q_W$ 
```

Model Free RL - Q-learning

So are we there yet ?

→ Further Problems with Q-learning

Model Free RL - Q-learning

So are we there yet ?

→ Further Problems with Q-learning

→ Although Q-update looks like gradient we have “moving targets”
(only partially addressed)

Model Free RL - Q-learning

So are we there yet ?

- Further Problems with Q-learning
 - Although Q-update looks like gradient we have “moving targets” (only partially addressed)
 - Convergence can be slow

Model Free RL - Q-learning

So are we there yet ?

- Further Problems with Q-learning
 - Although Q-update looks like gradient we have “moving targets” (only partially addressed)
 - Convergence can be slow
 - Reiterate: Efficient Exploration is key!

Model Free RL - Q-learning

So are we there yet ?

- Further Problems with Q-learning
 - Although Q-update looks like gradient we have “moving targets” (only partially addressed)
 - Convergence can be slow
 - Reiterate: Efficient Exploration is key!
 - What about continuous actions ?

Model Free RL - Q-learning

So are we there yet ?

- Further Problems with Q-learning
 - Although Q-update looks like gradient we have “moving targets” (only partially addressed)
 - Convergence can be slow
 - Reiterate: Efficient Exploration is key!
 - What about continuous actions ?
 - When are we really going to have a fully observed state ?